



Page Rank

Submitted By: Norina Akhtar
Student Id : 19643



Table Of Contents

[Introduction](#)

[Explanation](#)

[Manual Implementation / Design](#)

[implementation using Pyspark](#)

[PageRank + Scala + GCP](#)

[Enhancement Ideas](#)

[Conclusion](#)

Introduction

The original PageRank algorithm was described by Lawrence Page and Sergey Brin in several publications. It is given by

$$PR(A) = (1-d) + d (PR(T1)/C(T1) + \dots + PR(Tn)/C(Tn))$$

Where

- ❖ $PR(A)$ is the PageRank of page A,
- ❖ $PR(Ti)$ is the PageRank of pages Ti which link to page A,
- ❖ $C(Ti)$ is the number of outbound links on page Ti and
- ❖ d is a damping factor which can be set between 0 and 1.

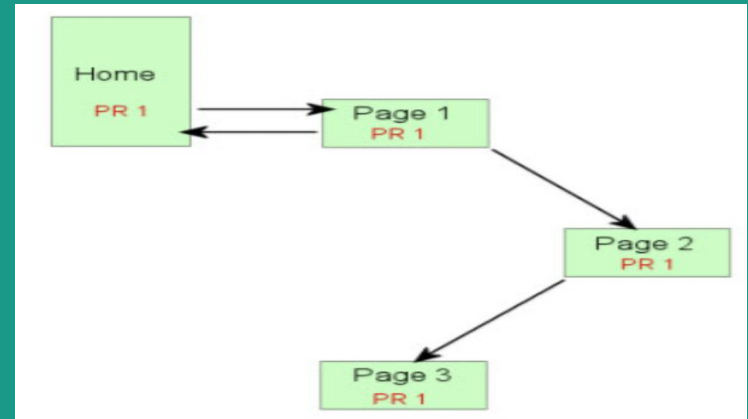
An iterative algorithm that performs many joins, so it is a good use case for RDD partitioning.

The algorithm maintains two datasets:

- $(pageID, linked\ List)$ elements containing the list of neighbors of each page,
- $(pageID, rank)$ elements containing the current rank for each page.

Explanation

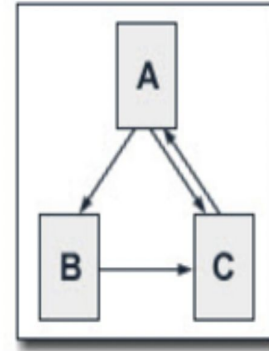
- We regard a small web consisting of three pages A, B, C, whereby page A links to the pages B and C, page B links to page C and page C links to Page A. According to Page and Brin, the damping factor d is usually set to 0.85.
- A web page does not have input will have
 - constant PageRank: $1-d$
 - the smallest PageRank
- Input Web Pages' impact to the PageRank of a web page
 - The more Input Web Pages the better.
 - The higher PageRank of an Input Web Page the better.



Manual Implementation / Design

Assuming

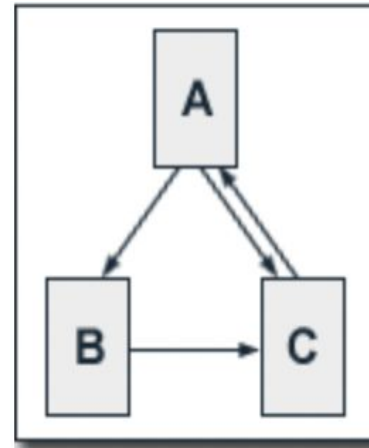
- the initial PageRank value for each webpage is 1.
- the damping factor is 0.85
- the relation of the webpages is:



Manual Implementation / Design

First Iteration:

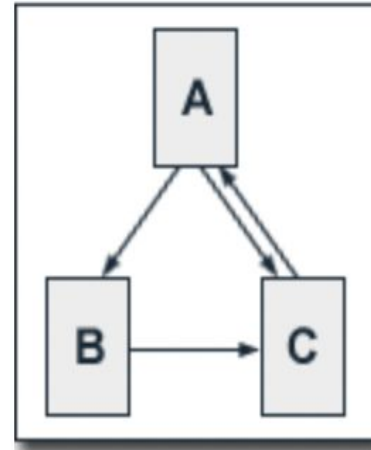
1. **PR(A)**
$$= (1-d) + d * (PR(C) / 1)$$
$$= (1-0.85) + 0.85 * (1)$$
$$= 1$$
2. **PR(B)**
$$= (1-d) + d * (PR(A) / 2)$$
$$= (1-0.85) + 0.85 * 0.5$$
$$= 0.575$$
3. **PR(C)**
$$= (1-d) + d * (PR(A) / 2 + PR(B) / 1)$$
$$= (1-0.85) + 0.85 * (0.5 + 1)$$
$$= 1.425$$



Manual Implementation

Second Iteration:

1. $PR(A)$
 $= 1 - 0.85 + 0.85 * 1.425$
 $= 1.36125$
2. $PR(B)$
 $= 1 - 0.85 + 0.85 * 0.5$
 $= 0.575$
3. $PR(C)$
 $1 - 0.85 + 0.85 * 1.075$
 $= 1.06375$



Implementation

Create cluster on GCP:

Start your Free Trial with \$300 in credit. Don't worry—you won't be charged if you run out of credits. [Learn more](#)

DISMISS

ACTIVATE

Google Cloud

CS570bigData

Search Products, resources, docs (/)

12N

Dataproc

Jobs on Clusters

Clusters

Jobs

Workflows

Autoscaling policies

Serverless

Batches

Metastore Services

Metastore

Federation

Utilities

Component exchange

Workbench

Create a Dataproc cluster on Compute Engine

Set up cluster

Begin by providing basic information.

Configure nodes (optional)

Change node compute and storage capabilities.

Customize cluster (optional)

Add cluster properties, features, and actions.

Manage security (optional)

Change access, encryption, and security settings.

CREATE

CANCEL

EQUIVALENT COMMAND LINE

Name

Cluster Name *
cluster-page-rank

Location

Region *
us-central1

Zone *
us-central1-f

Cluster type

☒ Standard (1 master, N workers)

☐ Single Node (1 master, 0 workers)
Provides one node that acts as both master and worker. Good for proof-of-concept or small-scale processing

☐ High Availability (3 masters, N workers)
Hadoop High Availability mode provides uninterrupted YARN and HDFS operations despite single-node failures or reboots

Autoscaling

Automates cluster resource management based on an autoscaling policy.



Implementation using Pyspark

Create new file:

vi pagerank.txt

```
Linux cluster-page-rank-m 3.10.0-116.el7.x86_64 #1 SMP Debian 3.10.117-2 (2022-07-28)
```

```
The programs included with the Debian GNU/Linux system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.
```

```
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent  
permitted by applicable law.
```

```
Last login: Mon Oct 31 23:33:02 2022 from 35.235.244.34
```

```
nakhtar@cluster-page-rank-m:~$ vi pagerank.txt
```

```
nakhtar@cluster-page-rank-m:~$ cat pagerank.txt
```

```
A B
```

```
A C
```

```
B C
```

```
C A
```



SSH-in-browser

A B

A C

B C

C A

~

~

~

~

~

~

~



Implementation

Create new directory on cluster i-e mydata:

```
hdfs dfs -mkdir hdfs:///mydata
```

```
nakhtar@cluster-page-rank-m:~$ hdfs dfs -mkdir hdfs:///mydata
```

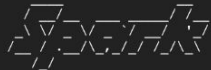
Put pagerank.txt file into hdfs directory:

```
nakhtar@cluster-page-rank-m:~$ hdfs dfs -put pagerank.txt hdfs:///mydata/pagerank.txt
```

```

nakhtar@cluster-page-rank-m:~$ pyspark
Python 3.8.13 | packaged by conda-forge | (default, Mar 25 2022, 06:04:10)
[GCC 10.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
22/11/01 05:32:14 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
22/11/01 05:32:14 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
22/11/01 05:32:14 INFO org.apache.spark.SparkEnv: Registering BlockManagerMasterHeartbeat
22/11/01 05:32:14 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
Welcome to

```



version 3.1.3

```

Using Python version 3.8.13 (default, Mar 25 2022 06:04:10)
Spark context Web UI available at http://cluster-page-rank-m.us-centrall-f.c.cs570bigdata.internal:41945
Spark context available as 'sc' (master = yarn, app id = application_1667279864037_0001).
SparkSession available as 'spark'.
>>> import re
>>> import sys
>>> from operator import add
>>>
>>> from pyspark.sql import SparkSession
>>> def computeContribs(urls, rank):
...     """Calculates URL contributions to the rank of other URLs."""
...     num_urls = len(urls)
...     for url in urls:
...         yield (url, rank / num_urls)
...
>>> def parseNeighbors(urls):
...     """Parses a urls pair string into urls pair."""
...     parts = re.split(r'\s+', urls)
...     return parts[0], parts[1]
...
>>> lines = spark.read.text("hdfs:///mydata/pagerank.txt").rdd.map(lambda r: r[0])
>>> lines.collect()
['A B', 'A C', 'B C', 'C A']
>>> links = lines.map(lambda urls: parseNeighbors(urls)).distinct().groupByKey().cache()
>>> links.collect()
[('A', <pyspark.resultiterable.ResultIterable object at 0x7f07c6413760>), ('B', <pyspark.resultiterable.ResultIterable object at 0x7f07c64137f0>), ('C', <pyspark.resultiterable.ResultIterable object at 0x7f07c6413850>)]
>>> ranks = links.map(lambda url_neighbors: (url_neighbors[0], 1.0))
>>> ranks.collect()
[('A', 1.0), ('B', 1.0), ('C', 1.0)]
>>> combine = links.join(ranks)
>>> combine.collect()
[('C', <pyspark.resultiterable.ResultIterable object at 0x7f07c641f8e0>, 1.0), ('A', <pyspark.resultiterable.ResultIterable object at 0x7f07c641f850>, 1.0), ('B', <pyspark.resultiterable.ResultIterable object at 0x7f07c641f910>, 1.0)]
>>> for iteration in range(int(10)):
...     contribs = combine.flatMap(lambda url_urls_rank: computeContribs(url_urls_rank[1][0], url_urls_rank[1][1]))
...     ranks = contribs.reduceByKey(lambda x,y:x+y)
...
>>> for (link, rank) in ranks.collect():
...     print("%s has rank: %s." % (link, rank))
...
C has rank: 1.5.
A has rank: 1.0.
B has rank: 0.5.
>>>

```



PageRank + Scala + GCP

Set up Scala on GCP:

- Create cluster
- Install scala using these commands

```
$ curl -fL https://github.com/coursier/launchers/raw/master/cs-x86_64-pc-linux.gz | gzip  
-d > cs && chmod +x cs && ./cs setup
```

```
$ export SCALA_HOME=/usr/local/share/scala
```

```
$ export PATH=$PATH:$SCALA_HOME/
```



PageRank + Scala + GCP

Create pagerank.txt file to store input data

```
Linux cluster-page-rank-m 3.10.0-1160.11.1.el7.x86_64 #1 SMP Debian 3.10.117-2 (2022-07-29)
```

```
The programs included with the Debian GNU/Linux system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.
```

```
Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent  
permitted by applicable law.
```

```
Last login: Mon Oct 31 23:33:02 2022 from 35.235.244.34
```

```
nakhtar@cluster-page-rank-m:~$ vi pagerank.txt
```

```
nakhtar@cluster-page-rank-m:~$ cat pagerank.txt
```

```
A B
```

```
A C
```

```
B C
```

```
C A
```



PageRank + Scala + GCP

Create new directory on cluster i-e mydata:

```
hdfs dfs -mkdir hdfs:///mydata
```

```
nakhtar@cluster-page-rank-m:~$ hdfs dfs -mkdir hdfs:///mydata
```

Put pagerank.txt file into hdfs directory:

```
hdfs dfs -put pagerank_data.txt hdfs:///mydata
```

```
nakhtar@cluster-page-rank-m:~$ hdfs dfs -put pagerank.txt hdfs:///mydata/pagerank.txt
```

Implementation using Scala


```
nakhtar@cluster-9aff-m:~$ curl -fL https://github.com/coursier/launchers/raw/master/cs-x86_64-pc-linux.gz | gzip -d > cs && chmod +x cs && ./cs setup
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
  0     0    0     0     0     0      0      0  --:--:-- --:--:-- --:--:--    0
100 20.0M 100 20.0M    0     0 15.3M    0  0:00:01  0:00:01 --:--:-- 33.6M

Checking if a JVM is installed
Found a JVM installed under /usr/lib/jvm/temurin-8-jdk-amd64.

Checking if ~/.local/share/coursier/bin is in PATH
Should we add ~/.local/share/coursier/bin to your PATH via ~/.profile? [Y/n] y

Checking if the standard Scala applications are installed
Installed ammonite
Installed cs
Installed coursier
Installed scala
Installed scalac
Installed scala-cli
Installed sbt
Installed sbt-n
Installed scalafmt

nakhtar@cluster-9aff-m:~$ export SCALA_HOME=/usr/local/share/scala
nakhtar@cluster-9aff-m:~$ export PATH=$PATH:$SCALA_HOME/
nakhtar@cluster-9aff-m:~$ vi pagerank.txt
nakhtar@cluster-9aff-m:~$ hdfs dfs -mkdir hdfs:///mydata
nakhtar@cluster-9aff-m:~$ hdfs dfs -put pagerank.txt hdfs:///mydata
nakhtar@cluster-9aff-m:~$ hdfs dfs -ls hdfs:///mydata
Found 1 items
-rw-r--r--  2 nakhtar hadoop      16 2022-11-01 20:58 hdfs:///mydata/pagerank.txt
nakhtar@cluster-9aff-m:~$ cat pagerank.txt
A B
A C
B C
C A
```

```
nakhtar@cluster-9aff-m:~$ spark-shell
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
22/11/01 20:59:47 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
22/11/01 20:59:47 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
22/11/01 20:59:47 INFO org.apache.spark.SparkEnv: Registering BlockManagerMasterHeartbeat
22/11/01 20:59:47 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
Spark context Web UI available at http://cluster-9aff-m.us-central1-c.c.cs570bigdata.internal:35735
Spark context available as 'sc' (master = yarn, app id = application_1667336002346_0001).
Spark session available as 'spark'.
Welcome to
```



```
Using Scala version 2.12.14 (OpenJDK 64-Bit Server VM, Java 1.8.0_345)
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala> val lines = sc.textFile("hdfs:///mydata/pagerank.txt")
lines: org.apache.spark.rdd.RDD[String] = hdfs:///mydata/pagerank.txt MapPartitionsRDD[1] at textFile at <console>:23

scala> val links = lines.map{ s => val parts = s.split("\\s+")
| (parts(0), parts(1))
| }.distinct().groupByKey().caching()
links: org.apache.spark.rdd.RDD[(String, Iterable[String])] = ShuffledRDD[11] at groupByKey at <console>:25

scala> var ranks = links.mapValues(v=> 1.0)
ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[12] at mapValues at <console>:23

scala> ranks.collect()
res1: Array[(String, Double)] = Array((B,1.0), (A,1.0), (C,1.0))
```



Result

1st Iteration:

```
scala> for (i <- 1 to 1){  
|   val contribs = links.join(ranks).values.flatMap(case (urls,rank) =>  
|   val size = urls.size  
|   urls.map(url => (url, rank/size))  
|   }  
|   ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)  
| }  
  
scala> val result = ranks.collect()  
result: Array[(String, Double)] = Array((B,0.575), (A,1.0), (C,1.4249999999999998))
```



Result

Second Iteration

```
val contribs = urls.map(url => (url, rank/size))
val size = urls.size
urls.map(url => (url, rank/size))
}
ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
}

scala> val result2 = ranks.collect()
result2: Array[(String, Double)] = Array((B,0.7285312499999999), (A,1.0541874999999998), (C,1.2172812499999996))

scala> result2.foreach(tup => println(tup._1 + "has rank: " + tup._2 + "."))
Bhas rank: 0.7285312499999999.
Ahas rank: 1.0541874999999998.
Chas rank: 1.2172812499999996.
```



Result

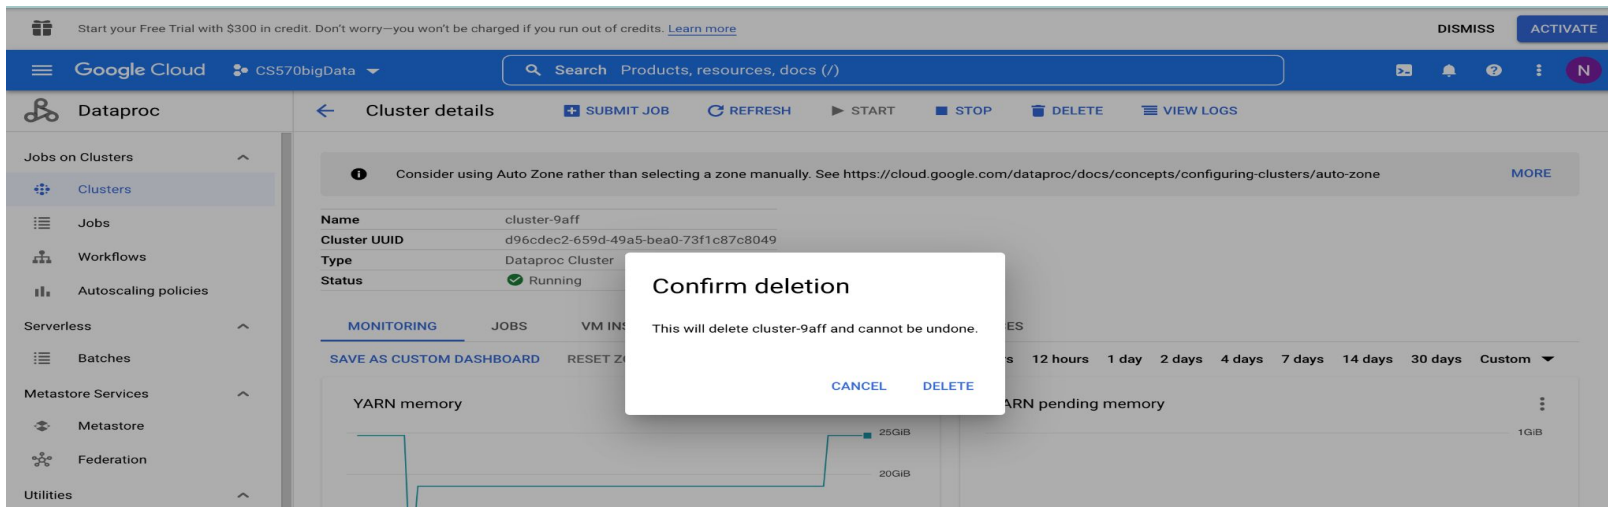
Third Iteration

```
scala> for (i <- 1 to 3){  
  |   val contribs = links.join(ranks).values.flatMap(case (urls,rank) =>  
  |   |   val size = urls.size  
  |   |   urls.map(url => (url, rank/size))  
  |   |  
  |   }  
  |   ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)  
  | }  
  
scala> val result3 = ranks.collect()  
result3: Array[(String, Double)] = Array((B,0.6534928515624998), (A,1.1375453730468745), (C,1.2089617753906245))  
  
scala> result3.foreach(tup => println(tup._1 + "has rank: " + tup._2 + "."))  
Bhas rank: 0.6534928515624998.  
Ahas rank: 1.1375453730468745.  
Chas rank: 1.2089617753906245.
```

So on.....

We can do it for n number iterations using for loop

Shut down the Cluster



The screenshot shows the Google Cloud Dataproc console interface. At the top, there's a header with the Google Cloud logo, account information (CS570bigData), and a search bar. Below the header, the left sidebar contains navigation options: Jobs on Clusters, Clusters (selected), Jobs, Workflows, Autoscaling policies, Serverless, Batches, Metastore Services, Metastore, Federation, and Utilities. The main content area displays the 'Cluster details' for a cluster named 'cluster-9aff'. The cluster is in a 'Running' status. A modal dialog titled 'Confirm deletion' is centered on the screen, with the text 'This will delete cluster-9aff and cannot be undone.' and two buttons: 'CANCEL' and 'DELETE'. The background is slightly dimmed to show the cluster details and monitoring graphs.

Start your Free Trial with \$300 in credit. Don't worry—you won't be charged if you run out of credits. [Learn more](#)

DISMISS ACTIVATE

Google Cloud CS570bigData

Search Products, resources, docs (/)

Dataproc

Cluster details

SUBMIT JOB REFRESH START STOP DELETE VIEW LOGS

Consider using Auto Zone rather than selecting a zone manually. See <https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/auto-zone> MORE

Name	cluster-9aff
Cluster UUID	d96cdec2-659d-49a5-bea0-73f1c87c8049
Type	Dataproc Cluster
Status	Running

MONITORING JOBS VM INSTANCES

SAVE AS CUSTOM DASHBOARD RESET ZONE

YARN memory

YARN pending memory

12 hours 1 day 2 days 4 days 7 days 14 days 30 days Custom

25GiB 20GiB 1GiB

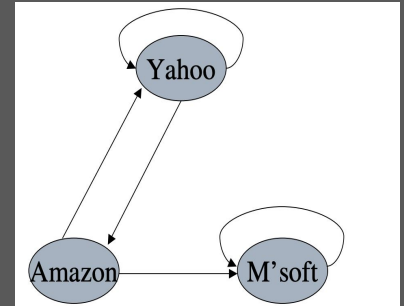
Confirm deletion

This will delete cluster-9aff and cannot be undone.

CANCEL DELETE

Enhancement Ideas

- We can calculate pagerank of big websites which have multiple links
- Compare scala and python performance
- We can test the the pagerank algorithm using n number of iteration





Conclusion

Hence pagerank is expressed as :

$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$

The PageRank value for a page u is dependent on the PageRank values for each page v contained in the set B_u (the set containing all pages linking to page u), divided by the number $L(v)$ of links from page v . The algorithm involves a damping factor(0.85) for the calculation of the PageRank. It is like the income tax which the govt extracts from one despite paying him itself



References

- [PageRank and design patterns for efficient graph algorithms](#)
- [Page Rank Tutorial](#)
- [A General Boosting Method and its Application to Learning Ranking Functions for Web](#)
- [Raise Your Google Ranking](#)